

de.NBI - Cloud Usermeeting 2021

Introduction to Kubernetes III:
RBAC/Security/Network policies

Sebastian Beyvers

`sebastian.beyvers@computational.bio.uni-giessen.de`

Marius Dieckmann

`marius.dieckmann@computational.bio.uni-giessen.de`

22 November 2021



Authorization in k8s

- Two types of accounts:
 - User
 - Serviceaccount
- Serviceaccounts are used in pods
 - Main use: Spawning new Pods
 - Can be mounted and used from programs in pods
- Authorization via RBAC (Role Based Access Control)
 - Defines "Roles" that describe permissions for resources
 - Roles are bound to accounts ("subjects")

Roles

- Roles define rights
- Two "Scopes":
 - Role: namespaced
 - ClusterRole: cluster scoped
- Uses verbs (get, list, create,...) and resources (pods, deployments,...) to define permission
- The assignment of a Role to a user/serviceaccounts is called Rolebinding (RB)
 - Rolebindings assign a role to collection of users, in a namespace
 - Cluster scoped rolebindings are called "ClusterRoleBindings"
 - ClusterRoles can be used in regular Rolebindings
 - Will only grant rights in the specified namespace

- Users can create/update roles
 - Requires the same permissions (namespaced)
 - Or requires the "escalate" permission (namespaced) [1.12+]
- Users can apply rolebindings
 - Requires the same permissions (namespaced)
 - Or requires the "bind" permission (namespaced)
- K8s has some predefined ClusterRoles (admin, edit, view)
 - Can be used as default
 - Can also be used by users for namespaced permissions

Pod Security Policy (PSP)

- Controls Pod security aspects ("permissions of pods")
- Cluster scoped resource
- Can be used in (Cluster)Roles like any other resource
- Pod security settings are validated against all available psp
- One of two important security measures in k8s (admin and user)

- Privileged
 - host device access
- Hostpath(s)
 - paths that can be mounted from the host
- Privilege escalation
 - uid/gid change permitted?
- RunAs[User,Group]
 - Forces the container to run as a specific user → avoid running container as root

- By default all Pods can communicate to each other
- Network policies restrict this network access
 - defaults are often applied by Admin
- Implemented in the network plugin
 - paths that can be mounted from the host
- Pod level firewall
- Introduces only very low latency
 - Sometimes in badly designed microservice application

Network policies II

- Rules are applied based on labels
 - e.g pod or/and namespace labels
- Rules can apply to:
 - Ingress
 - Egress
- Rules can be based on:
 - Labels
 - ipBlocks
 - DNS-Names
 - Ports
 - ...

Network policies – Best practices

- Use them
 - Check the default policy / behavior
- Be careful with label selectors
 - Namespace labels should only be handled by Admins
- Secure and limit access to all non-public APIs
 - gRPC
 - REST
 - Thrift
 - ...
- Avoid stacking multiple network policies
 - Hard to understand the behavior

Questions?